

Yash Maurya

maurya.yash152@gmail.com | yashmaurya.com | linkedin.com/in/yashmaurya | [Google Scholar](https://scholar.google.com/citations?user=...) | Pittsburgh, PA

EDUCATION

Carnegie Mellon University (CMU)

Master of Science in Information Technology - Privacy Engineering (MSIT-PE) | CGPA 3.9 / 4.0

Graduate Courses: *Federated Learning, Differential Privacy, Prompt Engineering, AI Governance, Responsible AI*

Research Areas: Unlearning in LLMs, Fairness, PETs(Privacy Enhancing Technologies), Synthetic Data, Implicit Bias Auditing

Pittsburgh, PA

Dec 2024

EXPERIENCE

Meta

Privacy Researcher

Pittsburgh, PA

Aug 2024 - Dec 2024

- Led privacy research examining consent flows' impact on AI engagement and data sharing across 500 participants
- Designed and evaluated 9 interface variants with different content lengths and privacy control granularity
- Found that medium/long content improved user comprehension, increased data transparency was desired, and interface length influenced privacy choices - longer interfaces encouraged basic settings while compact ones promoted granular control

Bank of New York Mellon (BNY)

AI Governance Intern

Pittsburgh, PA

June 2024 - Aug 2024

- Implemented LLM guardrails using Microsoft Presidio and NVIDIA NeMo for PII de-identification and content moderation
- Developed anti-jailbreaking rails using local LLMs like Llama-guard and ShieldGemma, enhancing AI system security and compliance
- Engineered automated pipelines to evaluate LLMs and RAG agents using benchmarks like MMLU, GSM8k, SALAD-Bench, and RAGAS
- Contributed to governance validation of Eliza, BNY's AI platform (15,000+ users), collaborating - Risk, Legal, Privacy, & Engineering teams

Samsung Electronics

R&D Engineer

Noida, India

July 2022 - Aug 2023

- Developed image narrative generation module using EfficientPS and UPSNet for panoptic segmentation in Samsung Discover 2.0
- Built large-scale data extraction, processing & ingestion engine for news articles using Selenium, BS4, processing 100k+ articles daily
- Engineered Unsupervised Topic Taxonomy construction pipeline for 10+ Million articles for Samsung News' recommendation system

DynamoFL (YC W22)

Federated Learning Researcher

San Francisco, CA | Remote

Feb 2021 - Aug 2021

- Implemented state-of-the-art Federated Learning(FL) algorithms from scratch including FedAvg, FedProx, FedMD, and FedHE
- Evaluated DP techniques like Laplacian and Gaussian noise algorithms using PyDP and prior-independent auctions for federated learning
- Engineered a PII sanitization portal leveraging Microsoft Presidio API and CTGAN for generating clean synthetic tabular data

CERTIFICATIONS

Certified Information Privacy Technologist (CIPT) | IAPP - International Association of Privacy Professionals | [Credential](https://iapp.org/certification/cipt/)

Jan 2024

PROJECTS

Guardrail Baselines for Unlearning in Large Language Models

Jan 2024 - May 2024

- Demonstrated that zero-shot prompting can achieve competitive unlearning performance on unlearning benchmarks without fine-tuning
- Extending the baseline by 16-bit/8-bit quantized fine-tuning LLaMA-2-7B using LoRA and QLoRA techniques for efficient unlearning
- Accepted at Secure and Trustworthy LLM(SetLLM) Workshop at **ICLR 2024**

UsersFirst: A User-Centric Privacy Threat Modeling Framework for Notice and Choice

Jan 2024 - May 2024

- Pioneered a novel threat modeling framework that addresses AI privacy vulnerabilities in user interface design
- Conducted in-depth interviews with 20 participants, validating framework efficacy vs. LINDDUN and PANOPTIC
- Integrated Privacy-by-Design principles to create actionable guidelines for combating deceptive design practices
- Accepted at Usable Privacy and Security (**SOUPS 2024**).

Unmasking Threats in Topics API (Replacement of Ad Cookies) | Presented at USENIX PEPR'24

Sept 2023 - Dec 2023

- Calculated Topics API's epsilon(privacy leakage budget) at 10.4 per week (epsilon > 10 signifies inadequate privacy protection)
- Our LLM based on Hierarchical BERT achieved 95.41% accuracy and 86.73% specificity for Membership Inference Attacks(MIA)
- Achieved 68.19% re-identification on an anonymized German Browsing Dataset, far surpassing Google's 1% claim

SELECTED PUBLICATIONS

Wang, T., Li, X. A., Rivera-Lanas, M., **Maurya, Y.**, Habib, H., Cranor, L. F., & Sadeh, N. "UsersFirst: A User-Centric Privacy Threat Modeling Framework for Notice and Choice" **SOUPS 2024**. https://www.usenix.org/system/files/soups2024_poster49_abstract-wang_final.pdf

P. Thaker, **Y. Maurya**, and V. Smith, "Guardrail Baselines for Unlearning in LLMs," **SET LLM@ICLR 2024**. <https://arxiv.org/abs/2403.03329>

Y. Maurya, P. Chandrahasan and P. G. "Federated Learning for Colorectal Cancer Prediction," 2022 **IEEE 3rd Global Conference for Advancement in Technology (GCAT)**, pp. 1-5, doi: [10.1109/GCAT55367.2022.9972224](https://doi.org/10.1109/GCAT55367.2022.9972224)

SKILLS

Programming Languages: Python, Java, C/C++, JavaScript, SQL, Rust, Bash

Libraries/Frameworks : PyTorch, TensorFlow, HuggingFace, OpenAI, Scikit-learn, Numpy, PySyft, Flower, Opacus, OpenDP, Nvidia NeMO

MLOps Tools & Frameworks: Wandb, Mlflow, Optuna, ZenML, Flask, Django, GCP, AWS, Docker, Langchain, W&B, Node.js, Neo4j, Airflow

Privacy Frameworks & Standards: NIST Privacy Framework, LINDDUN, MITRE PANOPTIC, FIPPs, OWASP, Privacy-by-Design, NIST AI RMF

Privacy Assessments & Documentation: Data Protection Impact Assessments(DPIAs), ROPAs, PIAs, Consent Management, Data Flow Mapping