Yash Maurya

maurya.yash152@gmail.com | Google Scholar | yashmaurya.com | yashmaurya | 4122142983

Professional Summary

Research engineer building **LLM evaluation and guardrails** across robustness, safety, and fairness. Post-training LLMs using **RLHF/DPO/GRPO** with **LLM-judged reward modeling** and automated eval loops.

Expertise: LLM Evaluation • Guardrails & Safety • Red-Teaming • Post-Training • Privacy • AI Governance
Skills: Python • PyTorch • Transformers • Hugging Face / LoRA • LLM Judges • W&B • Docker • AWS / GCP

Work Experience

Scale AI — San Francisco, CA
Research Engineer — *May 2025 – Present*

- Led **adversarial evaluation** and multi-agent + LLM-judge workflows for frontier models; built **AutoArchieQC** to reduce grading latency ~**840**× (1 week → 12 min) and conducted research on long-horizon agents and AI safety.

Bank of New York Mellon (BNY) — Pittsburgh, PA
Research Intern — *Jun 2024 – Aug 2024*

- Built an end-to-end **LLM eval + guardrails pipeline** (Mixtral/Llama-2/GPT-class) for accuracy, safety, and fairness with **real-time PII detection** (Presidio + NeMo) for enterprise deployment (**30k+** employees).

Samsung Electronics — Noida, India
Research Engineer — *Jul 2022 – Aug 2023*

- Scaled Samsung News content discovery with automated ingestion + unsupervised topic taxonomy (**100k+** articles/day; **10M+** items), improving recommendation quality and freshness.

DynamoFL (YC W22) — San Francisco, CA | Remote
Federated Learning Researcher — *Feb 2021 – Aug 2021*

- Built production-ready **federated learning + differential privacy workflows** (FedAvg/FedProx/FedMD/FedHE; PyDP) with **PII-aware synthetic data** (Presidio + CTGAN) to enable secure client demos.

Education

Carnegie Mellon University (CMU) — Pittsburgh, PA
M.S. Information Technology – Privacy Engineering (MSIT-PE), CGPA 3.9/4.0 — *Dec 2024*
Awards: IAPP Westin Scholar 2024; CMU Spark Entrepreneurship Grant Winner

Projects

Research Judge Training Environment — Jan 2025
Built an OpenEnv-compatible RL environment to train **LLM judges** for 2k-word research reports, with a multi-signal reward (accuracy + coverage + rationale) and synthetic corruption data generation.

Unmasking Threats in Google's Topics API — Sep 2023 – Dec 2023
Quantified privacy leakage ($\epsilon = 10.4$/week) and showed re-identification risk driven by edge cases / niche topics; explored **LLMs as adversarial attackers** (membership inference + re-ID) to stress-test differential privacy claims.

Selected Publications

[S = In Submission, C = Conference, W = Workshop]

- **[C.1] MoReBench: Evaluating Procedural and Pluralistic Moral Reasoning in LLMs** — *ICLR '26*
- **[S.1] LHAW: Controllable Underspecification for Long-Horizon Tasks** — *ICML '26 (sub.)*
- **[S.2] ROK-FORTRESS: Measuring the Effect of Geopolitical Transcreation for National Security and Public Safety** — *ICML '26 (sub.)*
- **[S.3] DualChem: Can LLMs Provide Dangerous Uplift in Dual Use Chemistry?** — *ICML '26 (sub.)*
- **[C.2] When Privacy Guarantees Meet Pre-Trained LLMs: A Case Study in Synthetic Data** — *USENIX PEPR '25*
- **[C.3] Position: LLM Unlearning Benchmarks are Weak Measures of Progress** — *SaTML '25*

Other venues of acceptance: SeT LLM@ICLR'24, USENIX SOUPS'22, RCV@CVPR'21, IEEE GCAT'21.