

Yash Maurya

✉ maurya.yash152@gmail.com 🎓 Google Scholar 🌐 yashmaurya.com in yashmaurya 📍 San Francisco

PROFESSIONAL SUMMARY

Research engineer building **LLM evaluation and guardrails** across agentic robustness, safety, and fairness. Post-training LLMs using **RLHF - PPO/DPO/GRPO** with **LLM-judged reward modeling**, rubric-based benchmarking, and automated eval loops.

Expertise: LLM Evaluation • Guardrails & Safety • Red-Teaming • Post-Training • Privacy • AI Governance

Skills: Python • PyTorch • Transformers • Hugging Face / LoRA • LLM Judges • W&B • Docker • AWS / GCP

WORK EXPERIENCE

Scale AI

Research Engineer

Pittsburgh, PA

May 2025 – Present

- Designed **rubric-based eval benchmarks** for subjective LLM capabilities across geopolitical safety (ROK-FORTRESS), dual-use chemistry uplift (DualChem), and pluralistic moral reasoning (MoReBench), using calibrated **LLM-as-judge panels** with expert-crafted criteria; and built a **synthetic underspecification pipeline** for stress-testing long-horizon agentic robustness (LHAW) (ICML '26, ICLR '26, In Submission)
- Built **AutoArchieQC** to reduce human-annotated data grading latency $\sim 840\times$ (1 week $\rightarrow$ 12 min); designed **human expert evaluation framework and experiments** for **Insights Generator**, a multi-agent system for corpus-level agent trace diagnostics with **30.4pp** scaffold improvement (In Submission).

Bank of New York Mellon (BNY)

Research Intern

Pittsburgh, PA

Jun 2024 – Aug 2024

- Built an end-to-end **LLM eval + guardrails pipeline** (Mixtral/Llama-2/GPT-class) for accuracy, safety, and fairness with **real-time PII detection** (Presidio + NeMo) for enterprise deployment (**30k+** employees).

Samsung Electronics

Research Engineer

Noida, India

Jul 2022 – Aug 2023

- Scaled Samsung News content discovery with automated ingestion + unsupervised topic taxonomy (**100k+** articles/day; **10M+** items), improving recommendation quality and freshness.

DynamoFL (YC W22)

Federated Learning Researcher

San Francisco, CA | Remote

Feb 2021 – Aug 2021

- Built production-ready **federated learning + differential privacy workflows** (FedAvg/FedProx/FedMD/FedHE; PyDP) with **PII-aware synthetic data** (Presidio + CTGAN) to enable secure client demos.

EDUCATION

Carnegie Mellon University (CMU)

M.S. Information Technology – Privacy Engineering (MSIT-PE), CGPA 3.9/4.0

Pittsburgh, PA

Dec 2024

Awards: IAPP Westin Scholar 2024; CMU Spark Entrepreneurship Grant Winner

SELECTED PUBLICATIONS

[S = In Submission, C = Conference, W = Workshop]

- [S.1] **Insights Generator: Systematic Corpus-Level Trace Diagnostics for LLM Agents** *In Submission*
- [W.1] **LHAW: Controllable Underspecification for Long-Horizon Tasks** *LLA@ICLR '26*
- [W.2] **ROK-FORTRESS: Measuring the Effect of Geopolitical Transcreation for National Security and Public Safety** *TAIGR@ICML '26*
- [S.2] **DualChem: Can LLMs Provide Dangerous Uplift in Dual Use Chemistry?** *In Submission*
- [C.1] **MoReBench: Evaluating Procedural and Pluralistic Moral Reasoning in LLMs** *ICLR '26*
- [C.2] **When Privacy Guarantees Meet Pre-Trained LLMs: A Case Study in Synthetic Data** *USENIX PEPR '25*
- [C.3] **Position: LLM Unlearning Benchmarks are Weak Measures of Progress** *SaTML '25*
- [C.4] **Through the Lens of LLMs: Unveiling Differential Privacy Challenges** *USENIX PEPR '24*
- [W.3] **Guardrail Baselines for Unlearning in LLMs** *SeT LLM@ICLR '24*

Other venues of acceptance: *SPP'26, USENIX SOUPS'22, RCV@CVPR'21, IEEE GCAT'21.*