

Yash Maurya

✉ maurya.yash152@gmail.com 🎓 Google Scholar 🌐 yashmaurya.com in yashmaurya 📍 San Francisco

PROFESSIONAL SUMMARY

Research engineer building **LLM evaluation and guardrails** across robustness, safety, and fairness. Post-training LLMs using **RLHF/DPO/GRPO** with **LLM-judged reward modeling** and automated eval loops.

Expertise: LLM Evaluation • Guardrails & Safety • Red-Teaming • Post-Training • Privacy • AI Governance

Skills: Python • PyTorch • Transformers • Hugging Face / LoRA • LLM Judges • W&B • Docker • AWS / GCP

WORK EXPERIENCE

Scale AI

Research Engineer

Pittsburgh, PA
May 2025 – Present

- Led multi-agent + LLM-judge evaluation for frontier models; built **AutoArchieQC** to reduce grading latency $\sim 840\times$ (1 week $\rightarrow$ 12 min) and developed research on long-horizon agents and AI safety (ICML sub., ICLR '26).

Bank of New York Mellon (BNY)

Research Intern

Pittsburgh, PA
Jun 2024 – Aug 2024

- Built an end-to-end **LLM eval + guardrails pipeline** (Mixtral/Llama-2/GPT-class) for accuracy, safety, and fairness with **real-time PII detection** (Presidio + NeMo) for enterprise deployment (**30k+** employees).

Samsung Electronics

Research Engineer

Noida, India
Jul 2022 – Aug 2023

- Scaled Samsung News content discovery with automated ingestion + unsupervised topic taxonomy (**100k+** articles/day; **10M+** items), improving recommendation quality and freshness.

DynamoFL (YC W22)

Federated Learning Researcher

San Francisco, CA | Remote
Feb 2021 – Aug 2021

- Built production-ready **federated learning + differential privacy workflows** (FedAvg/FedProx/FedMD/FedHE; PyDP) with **PII-aware synthetic data** (Presidio + CTGAN) to enable secure client demos.

EDUCATION

Carnegie Mellon University (CMU)

M.S. Information Technology – Privacy Engineering (MSIT-PE), CGPA 3.9/4.0

Pittsburgh, PA
Dec 2024

Awards: IAPP Westin Scholar 2024; CMU Spark Entrepreneurship Grant Winner

PROJECTS

Research Judge Training Environment

Jan 2025

Built an OpenEnv-compatible RL environment to train **LLM judges** for 2k-word research reports, with a multi-signal reward (accuracy + coverage + rationale) and synthetic corruption data generation.

Unmasking Threats in Google's Topics API

Sep 2023 – Dec 2023

Quantified privacy leakage ($\epsilon = 10.4/\text{week}$) and showed re-identification risk driven by edge cases / niche topics; explored **LLMs as adversarial attackers** (membership inference + re-ID) to stress-test differential privacy claims.

SELECTED PUBLICATIONS

[S = In Submission, C = Conference, W = Workshop]

- | | |
|---|-----------------|
| [C.1] MoReBench: Evaluating Procedural and Pluralistic Moral Reasoning in LLMs | ICLR '26 |
| [S.1] LHAW: Controllable Underspecification for Long-Horizon Tasks | ICML '26 (sub.) |
| [S.2] ROK-FORTRESS: Measuring the Effect of Geopolitical Transcreation for National Security and Public Safety | ICML '26 (sub.) |
| [S.3] DualChem: Can LLMs Provide Dangerous Uplift in Dual Use Chemistry? | ICML '26 (sub.) |
| [C.2] When Privacy Guarantees Meet Pre-Trained LLMs: A Case Study in Synthetic Data | USENIX PEPR '25 |
| [C.3] Position: LLM Unlearning Benchmarks are Weak Measures of Progress | SaTML '25 |

Other venues of acceptance: SeT LLM@ICLR'24, USENIX SOUPS'22, RCV@CVPR'21, IEEE GCAT'21.